

Kasuganti Venkata Mokshith Rao

✉ mokshithraokv@gmail.com ☎ 8522816334 📍 Hyderabad, India 🌐 LinkedIn 🏠 Github

CAREER OBJECTIVE

Data Science postgraduate with strong foundations in Python, machine learning, and software engineering fundamentals, along with hands-on experience building backend systems and REST API-based AI applications. Skilled in developing AI-powered applications using PyTorch, FastAPI, and Hugging Face, with exposure to NLP, LLM applications, and RAG systems. Seeking an entry-level technology role to apply my skills in software development, AI/ML, and backend engineering while contributing to real-world technology solutions.

EDUCATION

Master of Science in Data Science International Institute of Information Technology Hyderabad CGPA - 8.60	07/2024 – 07/2026 Hyderabad
Bachelor of Technology - CSE(AI & ML) CMR Institute of Technology CGPA - 7.70	12/2020 – 04/2024 Hyderabad

PROJECTS

AI Contract Intelligence System: Clause Detection & LLM-Based Rewriting 🔗

- Developed a full-stack AI application using **FastAPI** and **Streamlit** with REST API-based inference, contract upload, clause detection, and side-by-side clause rewriting.
- Fine-tuned **LegalBERT** using **LoRA** and **Hugging Face PEFT** on the CUAD dataset (~13K annotations) for contract clause detection across 41 categories, achieving ~70% accuracy and ~0.65 F1-score.
- Built a **Retrieval-Augmented Generation (RAG)** pipeline using the LEDGAR corpus, **FAISS**, and SentenceTransformers for context-aware clause retrieval.
- Integrated an **LLM-based rewriting module** to generate fair alternatives for potentially risky contract clauses.

SafeText AI: Cyberbullying Detection System 🔗

- Developed a full-stack NLP application using **FastAPI**, **Streamlit**, and **SQLite** for toxic-text classification and model comparison.
- Built and evaluated NLP classification approaches using **XGBoost**, **DistilBERT**, and **Llama-3.2-3B**, covering traditional ML, transformer-based, and LLM-based approaches.
- Achieved 94.3% accuracy and 0.97 ROC-AUC using XGBoost for binary toxic-text classification.
- Evaluated LLM-based classification, achieving **88.8% accuracy** and **87.9% recall** for toxic content detection.

SKILLS

- Programming & CS Fundamentals:** Python, SQL, OOP (Object-Oriented Programming), Data Structures & Algorithms (Basic), System Design Basics
- Generative AI & NLP:** RAG, LangChain, LLM Applications, Hugging Face, Transformers, SentenceTransformers, FAISS, Prompt Engineering, Agentic AI
- Machine Learning & Deep Learning:** PyTorch, Scikit-learn, Fine-tuning, LoRA/PEFT, Model Evaluation
- Backend Development & Databases:** FastAPI, REST APIs, Streamlit, MySQL, MongoDB, SQLite
- Data Processing & Tools:** Pandas, NumPy, Data Preprocessing, Git, GitHub, Jupyter Notebook, Google Colab, VS Code

CERTIFICATES

- RAG Systems Expert (Ready Tensor) - Sep, 2025
- AWS Academy Machine Learning Foundations - June, 2023
- Google Data Analytics (Coursera) - Feb, 2023
- Fundamentals of Deep Learning (Nvidia) - May, 2022